

Comparison between Sarsa and Expected Sarsa

Tom Bosc

10/02/17

Introduction

- ▶ TD methods combines model-free learning and bootstrapping.
- ▶ Sarsa: On-policy
- ▶ Q-learning: Off-policy
- ▶ Expected Sarsa: On-policy but also generalizes Q-learning

Update rules

Update rules:

- Sarsa:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$$

- Expected Sarsa:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \sum_{a'} \pi(a'|s') Q(s', a') - Q(s, a)]$$

Corresponding estimators (consider a, s, s' fixed here).

- Sarsa: $\hat{v}_t = r + \gamma Q(s', A')$
- Expected Sarsa: $v_t = r + \gamma \sum_{a'} \pi(a'|s') Q(s', a')$

Bias

Check that the 2 estimators are unbiased. Expectation over the policy (variable A').

$$E[v_t] = E[r + \gamma \sum_{a'} \pi(a'|s') Q(s', a')] = E[r] + \gamma E[V(s')] = E[r] + \gamma V(s')$$

$$E[\hat{v}_t] = E[r + \gamma Q(s', A')] = E[r] + \gamma V(s')$$

So $E[v_t] = E[\hat{v}_t]$.

Now, compute the variance.

Variance

Compute the variance:

$$\text{Var}(\hat{v}_t - v_t) = E[\hat{v}_t^2] - E[\hat{v}_t]^2 - (E[v_t^2] - E[v_t]^2)$$

$$\text{Var}(\hat{v}_t - v_t) = E[\hat{v}_t^2] - E[v_t^2]$$

$$\text{Var}(\hat{v}_t - v_t) = E[(r + \gamma Q(s', A'))^2] - E[(r + \gamma \sum_{a'} \pi(a'|s') Q(s', a'))^2]$$

$$\text{Var}(\hat{v}_t - v_t) = E[r] - E[r] + 2\gamma(E[Q(s', A')] - E[\sum_a Q(s', a)\pi(a|s')]) +$$

$$\gamma^2(E[Q(s', A')^2] - E[(\sum_a \pi(a|s') Q(s', a))^2])$$

$$\text{Var}(\hat{v}_t - v_t) = \gamma^2(E[Q(s', a')^2] - E[(\sum_a \pi(a|s') Q(s', a))^2])$$

Variance (2)

$$\text{Var}(\hat{v}_t - v_t) = \gamma^2 (E[Q(s', a')^2] - E[(\sum_a \pi(a|S') Q(S', a))^2])$$

$$\text{Var}(\hat{v}_t - v_t) = \gamma^2 (\sum_{a'} \pi(a|s') Q(s', a')^2) - E[(\sum_a \pi(a|s') Q(s', a))^2])$$

Variance (3)

The inner term is of the form:

$$\sum_i w_i x_i^2 - \left(\sum_i w_i x_i \right)^2, \quad (11)$$

where the w and x correspond to the π and Q values. When $w_i \geq 0$ for all i and $\sum_i w_i = 1$, we can give an unbiased estimate of the variance of the weighed values $w_i x_i$ as follows:

$$\frac{\sum_i w_i (x_i - \bar{x})^2}{1 - \sum_i w_i^2}, \quad (12)$$

where \bar{x} is the weighted mean $\sum_i w_i x_i$. Taking the numerator of this fraction and rewriting this gives us:



$$\begin{aligned} \sum_i w_i (x_i - \bar{x})^2 &= \sum_i w_i x_i^2 - 2 \sum_i w_i x_i \bar{x} + \sum_i w_i \bar{x}^2 \\ &= \sum_i w_i x_i^2 - 2 \bar{x}^2 + \bar{x}^2 \\ &= \sum_i w_i x_i^2 - \bar{x}^2, \end{aligned}$$

which is exactly the same quantity as given in (11). This

Experiment: Gambler's problem


Goal is to reach 32. $p_h = 0.3$

MC with exploring starts, 100k episodes.




Experiment: Gambler's problem

Sarsa, 300k updates, $\alpha = 0.1$, $\epsilon = 0.05$.



Expected Sarsa, 100k updates, $\alpha = 0.1$, $\epsilon = 0.05$.



Bibliography

- ▶ A Theoretical and Empirical Analysis of Expected Sarsa, Seijen et al., 2009